

SVM and KNN Classification

Question 1

I work on automated recognition. A typical binary classification problem is face recognition. Given two photos of face the system makes a determination as to whether these are two photos of the same person or different people.

Question 2 - 1

We will run a basic SVM using all the data. This is as per the notes on the assignment.

In [16]:

```
library("kernlab")
setwd("C:/Users/Jake/Documents/OMSA/Hw-1")
myData=as.matrix(read.table("credit_card_data-headers.txt",header=TRUE))
head(myData,5)

# fit model
model <- ksvm(myData[,1:10],myData[,11],type="C-svc",kernel="vanilladot",C=100,scaled=TRUE)

# calculate each and display the weights
a <- colSums(myData[model@SVindex,1:10]* model@coef[[1]])
a0 <- sum(a*myData[1,1:10]) - model@b

a
a0

pred <- predict(model,myData[,1:10])
pred

# display accuracy
sum(pred == myData[,11]) / nrow(myData)
```

| A1 | A2 | A3 | A8 | A9 | A10 | A11 | A12 | A14 | A15 | R1 |
|----|-------|-------|------|----|-----|-----|-----|-----|-----|----|
| 1 | 30.83 | 0.000 | 1.25 | 1 | 0 | 1 | 1 | 202 | 0 | 1 |
| 0 | 58.67 | 4.460 | 3.04 | 1 | 0 | 6 | 1 | 43 | 560 | 1 |
| 0 | 24.50 | 0.500 | 1.50 | 1 | 1 | 0 | 1 | 280 | 824 | 1 |
| 1 | 27.83 | 1.540 | 3.75 | 1 | 0 | 5 | 0 | 100 | 3 | 1 |
| 1 | 20.17 | 5.625 | 1.71 | 1 | 1 | 0 | 1 | 120 | 0 | 1 |

Setting default kernel parameters

A1

-0.000466036176627327

A2

-0.0140534983606244

A3

-0.0081688661743442

A8

0.0101292226736795

A9

0.501609468692229

A10

-0.00140343386065389

A11

0.00129121684002342

A12

-0.000266898857269382

A14

-0.206754961642446

A15

558.33559056503

-41.6013574057321

[illegible]

0.863914373088685

This model is clearly overfit given that we have used all the data. In the final question we will look to produce a model with a more reliable accuracy measure.

Question 2-2

We will fit a single KNN model. We will take it a step further and perform some training, comparing the best performance between 1 and 9.

In [9]:

```
#reload data
library("kkn")
myData <- read.table("credit_card_data-headers.txt",header=TRUE)
head(myData, 5)

for(i in 1:10) {
  myData[,i] <- as.numeric(as.character(Data[,i]))
}
myData$R1 <- as.factor(Data$R1)

# remove 50 entries for assesment of model
m <- dim(Data)[1]
Sample <- sample(1:m, 50)
testing <- myData[Sample, ]
learning <- myData[-Sample, ]

dim(learning)
dim(testing)

# assess models
model_train <- train.kknn(R1 ~ ., data=learning, scale=TRUE, kmax=9)
model_train

prediction <- predict(model_train, testing[, -11])
CM <- table(testing[, 11], prediction)
CM

accuracy <- (sum(diag(CM)))/sum(CM)
accuracy
```

| A1 | A2 | A3 | A8 | A9 | A10 | A11 | A12 | A14 | A15 | R1 |
|----|-------|-------|------|----|-----|-----|-----|-----|-----|----|
| 1 | 30.83 | 0.000 | 1.25 | 1 | 0 | 1 | 1 | 202 | 0 | 1 |
| 0 | 58.67 | 4.460 | 3.04 | 1 | 0 | 6 | 1 | 43 | 560 | 1 |
| 0 | 24.50 | 0.500 | 1.50 | 1 | 1 | 0 | 1 | 280 | 824 | 1 |
| 1 | 27.83 | 1.540 | 3.75 | 1 | 0 | 5 | 0 | 100 | 3 | 1 |
| 1 | 20.17 | 5.625 | 1.71 | 1 | 1 | 0 | 1 | 120 | 0 | 1 |

604 11

50 11

Call:

```
train.kknn(formula = R1 ~ ., data = learning, kmax = 9, scale = TRUE)
```

Type of response variable: nominal

Minimal misclassification: 0.1440397

Best kernel: optimal

Best k: 5

prediction

0 1

0 22 6

1 5 17

0.78

We achieved a 0.9 accuracy with $k=7$. We will use this to finalize a model.

In [5]:

```
# finalize model model
model <- kknn(R1 ~ ., train=learning, test=testing, scale=TRUE, k=7)
model

summary(model)
fit <- fitted(model)

CM <- table(testing$R1, fit)
CM

accuracy <- (sum(diag(CM)))/sum(CM)
accuracy
```

```
Call:
kknn(formula = R1 ~ ., train = learning, test = testing, k = 7,      scale
      = TRUE)
```

```
Response: "nominal"
```


Call:

```
kknn(formula = R1 ~ ., train = learning, test = testing, k = 7, scale  
= TRUE)
```

Response: "nominal"

| | fit | prob.0 | prob.1 |
|----|-----|------------|------------|
| 1 | 1 | 0.00000000 | 1.00000000 |
| 2 | 0 | 0.77080726 | 0.22919274 |
| 3 | 1 | 0.07275176 | 0.92724824 |
| 4 | 1 | 0.11329303 | 0.88670697 |
| 5 | 1 | 0.07275176 | 0.92724824 |
| 6 | 0 | 0.95945873 | 0.04054127 |
| 7 | 0 | 0.65944241 | 0.34055759 |
| 8 | 0 | 1.00000000 | 0.00000000 |
| 9 | 0 | 1.00000000 | 0.00000000 |
| 10 | 0 | 1.00000000 | 0.00000000 |
| 11 | 0 | 1.00000000 | 0.00000000 |
| 12 | 1 | 0.00000000 | 1.00000000 |
| 13 | 0 | 1.00000000 | 0.00000000 |
| 14 | 1 | 0.22919274 | 0.77080726 |
| 15 | 1 | 0.11136485 | 0.88863515 |
| 16 | 1 | 0.22919274 | 0.77080726 |
| 17 | 0 | 0.80313939 | 0.19686061 |
| 18 | 0 | 0.62686494 | 0.37313506 |
| 19 | 1 | 0.30194449 | 0.69805551 |
| 20 | 0 | 1.00000000 | 0.00000000 |
| 21 | 1 | 0.38946305 | 0.61053695 |
| 22 | 1 | 0.44588682 | 0.55411318 |
| 23 | 1 | 0.41367634 | 0.58632366 |
| 24 | 0 | 0.95945873 | 0.04054127 |
| 25 | 0 | 1.00000000 | 0.00000000 |
| 26 | 0 | 1.00000000 | 0.00000000 |
| 27 | 0 | 1.00000000 | 0.00000000 |
| 28 | 1 | 0.41367634 | 0.58632366 |
| 29 | 0 | 0.62686494 | 0.37313506 |
| 30 | 1 | 0.00000000 | 1.00000000 |
| 31 | 1 | 0.00000000 | 1.00000000 |
| 32 | 1 | 0.00000000 | 1.00000000 |
| 33 | 1 | 0.37313506 | 0.62686494 |
| 34 | 0 | 0.88863515 | 0.11136485 |
| 35 | 0 | 1.00000000 | 0.00000000 |
| 36 | 1 | 0.27163517 | 0.72836483 |
| 37 | 1 | 0.26973401 | 0.73026599 |
| 38 | 1 | 0.00000000 | 1.00000000 |
| 39 | 1 | 0.00000000 | 1.00000000 |
| 40 | 1 | 0.22919274 | 0.77080726 |
| 41 | 1 | 0.22919274 | 0.77080726 |
| 42 | 1 | 0.00000000 | 1.00000000 |
| 43 | 0 | 1.00000000 | 0.00000000 |
| 44 | 0 | 1.00000000 | 0.00000000 |
| 45 | 0 | 0.62686494 | 0.37313506 |
| 46 | 0 | 1.00000000 | 0.00000000 |
| 47 | 1 | 0.28247801 | 0.71752199 |
| 48 | 0 | 1.00000000 | 0.00000000 |
| 49 | 1 | 0.01274400 | 0.98725600 |
| 50 | 1 | 0.49724391 | 0.50275609 |

fit

0 1

0 23 5

1 0 22

0.9

Again we shouldn't expect this to generalise at 90%. In the next question we will use k-folds to produce more robust measures

Question 3

In [18]:

```
# I am going to use the caret library for classifiers
# detach("package:kernlab", unload=TRUE)
# detach("package:knn", unload=TRUE)
library(caret)

# Load data into df
myData <- read.table("credit_card_data-headers.txt", header=TRUE)

# convert to numeric data
for(i in 1:10) {
  myData[,i] <- as.numeric(as.character(myData[,i]))
}
myData$R1 <- as.factor(myData$R1)

# class distribution
cbind(freq=table(myData$R1), percentage=prop.table(table(myData$R1))*100)

# summarize correlations between input variables
complete_cases <- complete.cases(myData)
cor(myData[complete_cases,1:10])
```

| | freq | percentage |
|---|------|------------|
| 0 | 358 | 54.74006 |
| 1 | 296 | 45.25994 |

| | A1 | A2 | A3 | A8 | A9 | A10 |
|-----|-------------|-------------|-------------|-------------|-------------|-------------|
| A1 | 1.00000000 | 0.04580212 | -0.03475616 | 0.08433408 | -0.02227002 | 0.06603888 |
| A2 | 0.04580212 | 1.00000000 | 0.21573658 | 0.40910052 | 0.22151088 | -0.09674680 |
| A3 | -0.03475616 | 0.21573658 | 1.00000000 | 0.30043424 | 0.23678042 | -0.16580972 |
| A8 | 0.08433408 | 0.40910052 | 0.30043424 | 1.00000000 | 0.33513714 | -0.22853844 |
| A9 | -0.02227002 | 0.22151088 | 0.23678042 | 0.33513714 | 1.00000000 | -0.42878000 |
| A10 | 0.06603888 | -0.09674680 | -0.16580972 | -0.22853844 | -0.42878000 | 1.00000000 |
| A11 | -0.01719151 | 0.19245630 | 0.26967523 | 0.32758743 | 0.37722077 | -0.56940610 |
| A12 | -0.05131107 | -0.05158808 | 0.00626837 | -0.13991523 | -0.08842336 | 0.02145246 |
| A14 | 0.07642180 | -0.09013265 | -0.21710129 | -0.06388272 | -0.05866532 | 0.03748592 |
| A15 | 0.01296216 | 0.02790891 | 0.11972427 | 0.05224694 | 0.08418701 | -0.06832104 |

In [22]:

```
set.seed(42)
validation_index <- createDataPartition(myData$R1, p=0.80, list=FALSE)
validation <- myData[-validation_index,]
temp <- myData[validation_index,]
test_index <- createDataPartition(temp$R1, p=0.75, list=FALSE)
testing <- temp[-test_index,]
training <- temp[test_index,]

# manually running c through ksvm, I found insensitivity with the vanilladot kernel and
# with rbf diminished
# as I moved away from 1.
x <- training[,1:10]
x <- as.matrix(x)
y <- as.numeric(training$R1)
x_test <- testing[,1:10]
x_test <- as.matrix(x_test)
y_test <- as.numeric(testing$R1)
c=seq(1,200, by=20)
for(i in 1:10) {
  model<-ksvm(x,y,type="C-svc",kernel="rbf",C=c[i],scaled=TRUE)
  pred_svm <- predict(model,x_test)
  CM <- table(y_test, pred_svm)
  accuracy <- (sum(diag(CM)))/sum(CM)
  print(c[i])
  print(CM)
  print(accuracy)
}

# Convert DF values to numeric
for(i in 1:10) {
  myData[,i] <- as.numeric(as.character(myData[,i]))
}
myData$R1 <- as.factor(myData$R1)

# manually running k through kknn
k=seq(1,20,by=1)
for(i in 1:20) {
  model_knn <- kknn(R1 ~ ., train=training, test=testing, scale=TRUE, k=k[i])
  fit <- fitted(model_knn)
  CM <- table(testing$R1, fit)
  print(k[i])
  print(CM)
  accuracy <- (sum(diag(CM)))/sum(CM)
  print(accuracy)
}
```

```

[1] 1
      pred_svm
y_test 1 2
      1 58 13
      2 6 53
[1] 0.8538462
[1] 21
      pred_svm
y_test 1 2
      1 62 9
      2 16 43
[1] 0.8076923
[1] 41
      pred_svm
y_test 1 2
      1 61 10
      2 15 44
[1] 0.8076923
[1] 61
      pred_svm
y_test 1 2
      1 61 10
      2 20 39
[1] 0.7692308
[1] 81
      pred_svm
y_test 1 2
      1 61 10
      2 19 40
[1] 0.7769231
[1] 101
      pred_svm
y_test 1 2
      1 61 10
      2 20 39
[1] 0.7692308
[1] 121
      pred_svm
y_test 1 2
      1 61 10
      2 20 39
[1] 0.7692308
[1] 141
      pred_svm
y_test 1 2
      1 60 11
      2 20 39
[1] 0.7615385
[1] 161
      pred_svm
y_test 1 2
      1 60 11
      2 20 39
[1] 0.7615385
[1] 181
      pred_svm
y_test 1 2
      1 60 11
      2 21 38
[1] 0.7538462
[1] 1

```

```
fit
  0  1
0 61 10
1 17 42
[1] 0.7923077
[1] 2
fit
  0  1
0 61 10
1 17 42
[1] 0.7923077
[1] 3
fit
  0  1
0 61 10
1 17 42
[1] 0.7923077
[1] 4
fit
  0  1
0 61 10
1 17 42
[1] 0.7923077
[1] 5
fit
  0  1
0 62  9
1 16 43
[1] 0.8076923
[1] 6
fit
  0  1
0 63  8
1 16 43
[1] 0.8153846
[1] 7
fit
  0  1
0 63  8
1 17 42
[1] 0.8076923
[1] 8
fit
  0  1
0 63  8
1 17 42
[1] 0.8076923
[1] 9
fit
  0  1
0 63  8
1 18 41
[1] 0.8
[1] 10
fit
  0  1
0 62  9
1 15 44
[1] 0.8153846
[1] 11
fit
```

```
      0  1
    0 62  9
    1 15 44
[1] 0.8153846
[1] 12
    fit
      0  1
    0 62  9
    1 15 44
[1] 0.8153846
[1] 13
    fit
      0  1
    0 61 10
    1 14 45
[1] 0.8153846
[1] 14
    fit
      0  1
    0 61 10
    1 12 47
[1] 0.8307692
[1] 15
    fit
      0  1
    0 61 10
    1 12 47
[1] 0.8307692
[1] 16
    fit
      0  1
    0 61 10
    1 12 47
[1] 0.8307692
[1] 17
    fit
      0  1
    0 61 10
    1 13 46
[1] 0.8230769
[1] 18
    fit
      0  1
    0 61 10
    1 13 46
[1] 0.8230769
[1] 19
    fit
      0  1
    0 61 10
    1 13 46
[1] 0.8230769
[1] 20
    fit
      0  1
    0 61 10
    1 14 45
[1] 0.8153846
```

The classes are marginally unbalanced, the correlation looks good. I am not going to remove any data.

The approach before used $k = 10$ for k-folds cross validation with 3 repeats. We then run it through a grid search to find the best performing data. Following this we will assess the model accuracy on the validation data.

In [23]:

```
# Tune SVM
# 10-fold cross validation with 3 repeats
control <- trainControl(method="repeatedcv", number=10, repeats=3)
metric <- "Accuracy"
set.seed(7)
grid <- expand.grid(.sigma=c(0.025, 0.05, 0.1, 0.15), .C=seq(1, 200, by=20))
fit.svm <- train(R1~., data=training, method="svmRadial", metric=metric, tuneGrid=grid,
  preProc=c("scale"), trControl=control)
print(fit.svm)
plot(fit.svm)

# Tune kNN
# 10-fold cross validation with 3 repeats
control <- trainControl(method="repeatedcv", number=10, repeats=3)
metric <- "Accuracy"
set.seed(7)
grid <- expand.grid(.k=seq(1,20,by=1))
fit.knn <- train(R1~., data=training, method="knn", metric=metric, tuneGrid=grid, prePr
oc=c("scale"), trControl=control)
print(fit.knn)
plot(fit.knn)
```


Support Vector Machines with Radial Basis Function Kernel

394 samples
10 predictor
2 classes: '0', '1'

Pre-processing: scaled (10)

Resampling: Cross-Validated (10 fold, repeated 3 times)

Summary of sample sizes: 354, 354, 355, 355, 355, ...

Resampling results across tuning parameters:

| sigma | C | Accuracy | Kappa |
|-------|-----|-----------|-----------|
| 0.025 | 1 | 0.8729341 | 0.7484916 |
| 0.025 | 21 | 0.8687910 | 0.7389673 |
| 0.025 | 41 | 0.8596457 | 0.7194125 |
| 0.025 | 61 | 0.8512045 | 0.7015536 |
| 0.025 | 81 | 0.8477215 | 0.6940734 |
| 0.025 | 101 | 0.8494962 | 0.6976564 |
| 0.025 | 121 | 0.8469523 | 0.6924786 |
| 0.025 | 141 | 0.8452429 | 0.6889450 |
| 0.025 | 161 | 0.8452002 | 0.6888132 |
| 0.025 | 181 | 0.8468882 | 0.6920568 |
| 0.050 | 1 | 0.8712461 | 0.7450740 |
| 0.050 | 21 | 0.8485751 | 0.6959329 |
| 0.050 | 41 | 0.8494309 | 0.6974511 |
| 0.050 | 61 | 0.8527429 | 0.7036564 |
| 0.050 | 81 | 0.8493893 | 0.6965780 |
| 0.050 | 101 | 0.8443252 | 0.6858923 |
| 0.050 | 121 | 0.8357973 | 0.6686327 |
| 0.050 | 141 | 0.8307760 | 0.6583348 |
| 0.050 | 161 | 0.8274438 | 0.6512451 |
| 0.050 | 181 | 0.8265891 | 0.6492818 |
| 0.100 | 1 | 0.8653261 | 0.7331355 |
| 0.100 | 21 | 0.8476361 | 0.6925564 |
| 0.100 | 41 | 0.8374629 | 0.6715271 |
| 0.100 | 61 | 0.8382535 | 0.6723041 |
| 0.100 | 81 | 0.8383187 | 0.6724730 |
| 0.100 | 101 | 0.8256455 | 0.6466661 |
| 0.100 | 121 | 0.8205601 | 0.6364893 |
| 0.100 | 141 | 0.8078464 | 0.6101618 |
| 0.100 | 161 | 0.8010717 | 0.5967633 |
| 0.100 | 181 | 0.8027598 | 0.6004378 |
| 0.150 | 1 | 0.8595153 | 0.7211370 |
| 0.150 | 21 | 0.8408176 | 0.6782704 |
| 0.150 | 41 | 0.8315879 | 0.6592319 |
| 0.150 | 61 | 0.8171413 | 0.6305625 |
| 0.150 | 81 | 0.8035279 | 0.6025607 |
| 0.150 | 101 | 0.7950900 | 0.5853625 |
| 0.150 | 121 | 0.7916700 | 0.5782091 |
| 0.150 | 141 | 0.7900259 | 0.5741907 |
| 0.150 | 161 | 0.7857951 | 0.5649980 |
| 0.150 | 181 | 0.7875270 | 0.5688254 |

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were sigma = 0.025 and C = 1.

k-Nearest Neighbors

394 samples
10 predictor
2 classes: '0', '1'

Pre-processing: scaled (10)


Resampling: Cross-Validated (10 fold, repeated 3 times)


Summary of sample sizes: 354, 354, 355, 355, 355, ...

Resampling results across tuning parameters:

| k | Accuracy | Kappa |
|----|-----------|-----------|
| 1 | 0.8246188 | 0.6449127 |
| 2 | 0.8037213 | 0.6025591 |
| 3 | 0.8392578 | 0.6761542 |
| 4 | 0.8374640 | 0.6727079 |
| 5 | 0.8400270 | 0.6790351 |
| 6 | 0.8469984 | 0.6925101 |
| 7 | 0.8504161 | 0.6995952 |
| 8 | 0.8486640 | 0.6961528 |
| 9 | 0.8546064 | 0.7075619 |
| 10 | 0.8495625 | 0.6973627 |
| 11 | 0.8522143 | 0.7017616 |
| 12 | 0.8530061 | 0.7032534 |
| 13 | 0.8437101 | 0.6839648 |
| 14 | 0.8437303 | 0.6839817 |
| 15 | 0.8446278 | 0.6850539 |
| 16 | 0.8479791 | 0.6917892 |
| 17 | 0.8572323 | 0.7103834 |
| 18 | 0.8546907 | 0.7048272 |
| 19 | 0.8615081 | 0.7187452 |
| 20 | 0.8565092 | 0.7086244 |

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 19.





We found $k=19$ to be best with 86% accuracy. SVM experienced similar performance with $\sigma = 0.025$ and $C = 1$ I will build each model we can compare each against the validation data.

In [24]:

```
# build final model svm
x <- training[,1:10]
x <- as.matrix(x)
y <- as.numeric(training$R1)
model_svm <- ksvm(x, y,type="C-svc",kernel="rbf",simga=0.025,C=1)

x_valid <- as.matrix(validation[,1:10])
y_valid <- as.numeric(validation$R1)
pred_svm <- predict(model,x_valid)
CM <- table(y_valid, pred_svm)
CM
accuracy <- (sum(diag(CM)))/sum(CM)
accuracy

# Acheived 83.8 accuracy on the validation set
```

```
      pred_svm
y_valid 1  2
      1 59 12
      2 13 46
```

0.807692307692308

In [25]:

```
# build final model knn
model_knn <- kknn(R1 ~ ., train=training, test=validation, scale=TRUE, k=19)
fit <- fitted(model_knn)
CM <- table(validation$R1, fit)
CM
accuracy <- (sum(diag(CM)))/sum(CM)
accuracy
```

```
      fit
      0  1
0 56 15
1  8 51
```

0.823076923076923

The accuracy on the validation suggest our accuracy with the algorithms is 80.7% for the svm and 82.3% for knn.