

Hw-4

Homework 4

Q1

Below we import the data, perform principle component analysis and train a basic linear model with the first four components (highest variance explained).

```
# load the libraries
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
crime = read.table('../Hw-3/uscrime.txt', header = TRUE)
```

```
crime.pca <- prcomp(crime[1:15], scale. = TRUE)
```

```
summary(crime.pca)
```


```
## Importance of components%s:
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  2.4534 1.6739 1.4160 1.07806 0.97893 0.74377
## Proportion of Variance 0.4013 0.1868 0.1337 0.07748 0.06389 0.03688
## Cumulative Proportion 0.4013 0.5880 0.7217 0.79920 0.86308 0.89996
##          PC7      PC8      PC9      PC10     PC11     PC12
## Standard deviation  0.56729 0.55444 0.48493 0.44708 0.41915 0.35804
## Proportion of Variance 0.02145 0.02049 0.01568 0.01333 0.01171 0.00855
## Cumulative Proportion 0.92142 0.94191 0.95759 0.97091 0.98263 0.99117
##          PC13     PC14     PC15
## Standard deviation  0.26333 0.2418 0.06793
## Proportion of Variance 0.00462 0.0039 0.00031
## Cumulative Proportion 0.99579 0.9997 1.00000
```


The first 4 elements explain 80% of the variance. Lets have a look at the breakdown graphically:

```
plot(crime.pca, type = "l")
```

crime.pca



```
biplot(crime.pca, scale = 0)
```



From the elbow graph we see the diminishing return each component contributes. From the biplot we see the rotation and level of importance.

```
sd <- crime.pca$sdev
loadings <- crime.pca$rotation
rownames(loadings) <- colnames(crime[1:15])
scores <- crime.pca$x

crime.train <- as.data.frame(crime.pca$x[,1:4], header = T)
crime.train$Y <- crime$Crime

model.pca <- lm(Y~., data=crime.train)
summary(model.pca)
```

```
##
## Call:
## lm(formula = Y ~ ., data = crime.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -557.76 -210.91  -29.08  197.26  810.35
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   905.09      49.07  18.443 < 2e-16 ***
## PC1             65.22      20.22   3.225  0.00244 **
## PC2            -70.08      29.63  -2.365  0.02273 *
```

```
## PC3          25.19      35.03   0.719  0.47602
## PC4          69.45      46.01   1.509  0.13872
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 336.4 on 42 degrees of freedom
## Multiple R-squared:  0.3091, Adjusted R-squared:  0.2433
## F-statistic: 4.698 on 4 and 42 DF,  p-value: 0.003178
```

With an r^2 of 30.1% we explain less variance than last weeks homework. Below we can explore preprocessing the parameters in order to reverse engineer the model back into the original factors. I will use 10-fold cross validation in the training.

```
# calculate the pre-process parameters from the dataset
preprocessParams <- preProcess(crime[1:15], method=c("center", "scale", "pca"), pcaComp = 4)
# preprocessParams$rotation

# transform the dataset using the parameters
transformed <- predict(preprocessParams, crime[1:15])
# summarize the transformed dataset
# summary(transformed)
transformed$Y <- crime$Crime

control <- trainControl(method="repeatedcv", number=10, repeats=3)
metric <- "RMSE"
# lm
set.seed(7)
fit.lm <- train(Y~., data=transformed, method="lm", metric=metric, trControl=control)
final.model <- fit.lm$finalModel

model.inter<-as.data.frame(preprocessParams$rotation)
for (i in 1:4){
  model.inter[,i] <- model.inter[,i]*final.model$coefficients[i+1]
}

coefs <- data.frame(row.names = c(colnames(crime[1:15]), "intercept"))
for (i in 1:15){
  coefs[i,1] <- sum(model.inter[i,])
}
coefs[16,1]<-final.model$coefficients[1]
coefs
```

```
##          V1
## M      -21.277963
## So      10.223091
## Ed      14.352610
## Po1     63.456426
## Po2     64.557974
## LF     -14.005349
## M.F    -24.437572
## Pop     39.830667
## NW      15.434545
## U1     -27.222281
## U2       1.425902
## Wealth  38.607855
```

```
## Ineq      -27.536348
## Prob      3.295707
## Time      -6.612616
## intercept 905.085106
```

So this is our model translated back into the original elements in via “un-rotation”

Q2

Below we build the models for the regression tree and the random forest:

```
library(caret)
library(rpart)
library(corrplot)

crime = read.table('../Hw-3/uscrime.txt', header = TRUE)

# Run algorithms using 10-fold cross validation
control <- trainControl(method="repeatedcv", number=10, repeats=3)
metric <- "RMSE"

set.seed(42)
grid <- expand.grid(.cp=c(0, 0.05, 0.1))
fit.cart <- train(Crime~., data=crime, method="rpart", metric=metric, tuneGrid=grid, preProc=c("center"))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
## trainInfo, : There were missing values in resampled performance measures.


# Random Forest
set.seed(42)
fit.rf <- train(Crime~., data=crime, method="rf", metric=metric, preProc=c("BoxCox"), trControl=control)

## Loading required package: randomForest
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
##
# Compare algorithms
transform_results <- resamples(list(CART=fit.cart, RF=fit.rf))
summary(transform_results)

##
## Call:
## summary.resamples(object = transform_results)
##
## Models: CART, RF
## Number of resamples: 30
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
```

```
## CART 143.56524 247.7585 342.3559 368.2294 467.7376 682.7395 0
## RF 75.71638 176.4595 260.3188 272.0210 346.3054 537.0275 0
##
## Rsquared
## Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## CART 0.0001149384 0.1016129 0.3311003 0.3735410 0.6683652 0.9818932 3
## RF 0.0008630734 0.4658528 0.7460322 0.6717879 0.9498149 0.9901082 0
```

```
dotplot(transform_results)
```



The RF model had a slightly lower RMSE, They had very high R-sq values. Lets have a look at the look at each of these to understand the results:

```
par(mfrow = c(1,2), xpd = NA)
plot(fit.cart$finalModel, uniform=TRUE,
     main="Regression Tree for USCrime")
text(fit.cart$finalModel, use.n=TRUE, all=TRUE, cex=.8)

print(fit.rf$finalModel) # view results
```

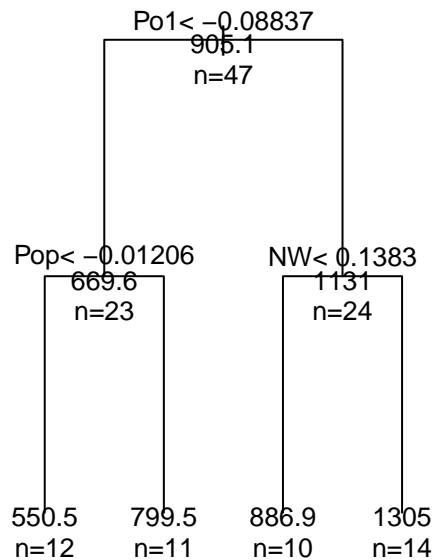
```
##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 8
##
##           Mean of squared residuals: 88026.1
```

```
## % Var explained: 39.87
```

```
importance(fit.rf$finalModel, type=2) # importance of each predictor
```

```
##      IncNodePurity
## M      247503.28
## So      18299.65
## Ed      197872.76
## Po1     1296476.21
## Po2     1284771.18
## LF      201132.67
## M.F     244774.20
## Pop     274764.11
## NW      494086.68
## U1      111766.71
## U2      152780.11
## Wealth  718980.85
## Ineq    212007.22
## Prob    912158.35
## Time    157958.49
```

Regression Tree for USCrime



We can see that there were quite a lot of impurity in certain nodes. As with previous analysis we can probably prune down some of the features to get a higher performing model. From the tree structure, we can see that the decisions are `Po1`, `Pop` and `NW`.

Q3

We use logistic regression nodes in deep learning algorithms in combination with relus to do certain computer vision tasks. So looking at a sequence of pixels it might say this image is a possible case of fraud. Another example may be when considering which sites to drill for oil. Features may include: mineral sample levels, location of previous oil sites etc.

Q4

Lets produce our logistic regression model


```
credit = read.table("germancredit.txt")

library(ggplot2)

names(credit) <- c('ca_status','mob','credit_history','purpose','credit_amount','savings',
                  'present_employment_since','status_sex','installment_rate_income','other_debtors',
                  'present_residence_since','property','age','other_installment','housing','existing_credits',
                  'job','liable_maintenance_people','telephone','foreign_worker','gb')

credit$gb <- factor(credit$gb,levels=c(2,1),labels=c("bad","good"))
model <- glm(data=credit,formula=gb~.,family=binomial(link="logit"))
credit$prob <- predict(model,newdata=credit, type="response")

ggplot(data=credit) + geom_boxplot(aes(y=prob,x=gb)) + coord_flip()
```



Just from looking at the distribution the model is doing a reasonable job of separating the data with more

than 75% of the good credit ratings above 75% of the bad credit ratings.

Now we will have a look at the ROC and AUC measures:

```
library(ROCR)
```

```
## Loading required package: gplots
```

```
##
```

```
## Attaching package: 'gplots'
```

```
## The following object is masked from 'package:stats':
```


```
##
```

```
##      lowess
```

```
pr <- prediction(credit$prob, credit$gb)
```

```
prf <- performance(pr, measure = "tpr", x.measure = "fpr")
```

```
plot(prf)
```



```
auc <- performance(pr, measure = "auc")
```

```
auc <- auc@y.values[[1]]
```

```
auc
```

```
## [1] 0.833781
```

this looks pretty good, we see a good climb in performance, and an auc of 83% certainly suggests this may be a useful model.

Finally, lets go about thresholding this. We will define a cost function and then just descend this from 50% through to where cost starts to increase:

```

thresh <- 0.5
cost <- Inf

repeat{
  fitted.results <- ifelse(credit$prob > thresh,"good","bad")
  CM <- table(credit$gb, fitted.results)
  new_cost = 5*CM[1,2]+CM[2,1]
  if(new_cost < cost){
    cost <- new_cost
    thresh <- thresh +0.01
  }
  else{
    thresh <- thresh - 0.01
    break
  }
}
thresh

```

```
## [1] 0.69
```

We produce a threshold of 69% (lol), lets have a look at what the model performance is (on the training data) at this threshold:

```

fitted.results <- ifelse(credit$prob > thresh,"good","bad")
CM <- table(credit$gb, fitted.results)
thresh

```

```
## [1] 0.69
```

```
CM
```

```

##      fitted.results
##      bad good
## bad  231   69
## good 167  533

```

```

misClasificError <- mean(fitted.results != credit$gb)
print(paste('Accuracy',1-misClasificError))

```

```
## [1] "Accuracy 0.764"
```

76.4% accuracy however this largely consists of good results classified as bad, rather than bad classified as good.

What a great time.