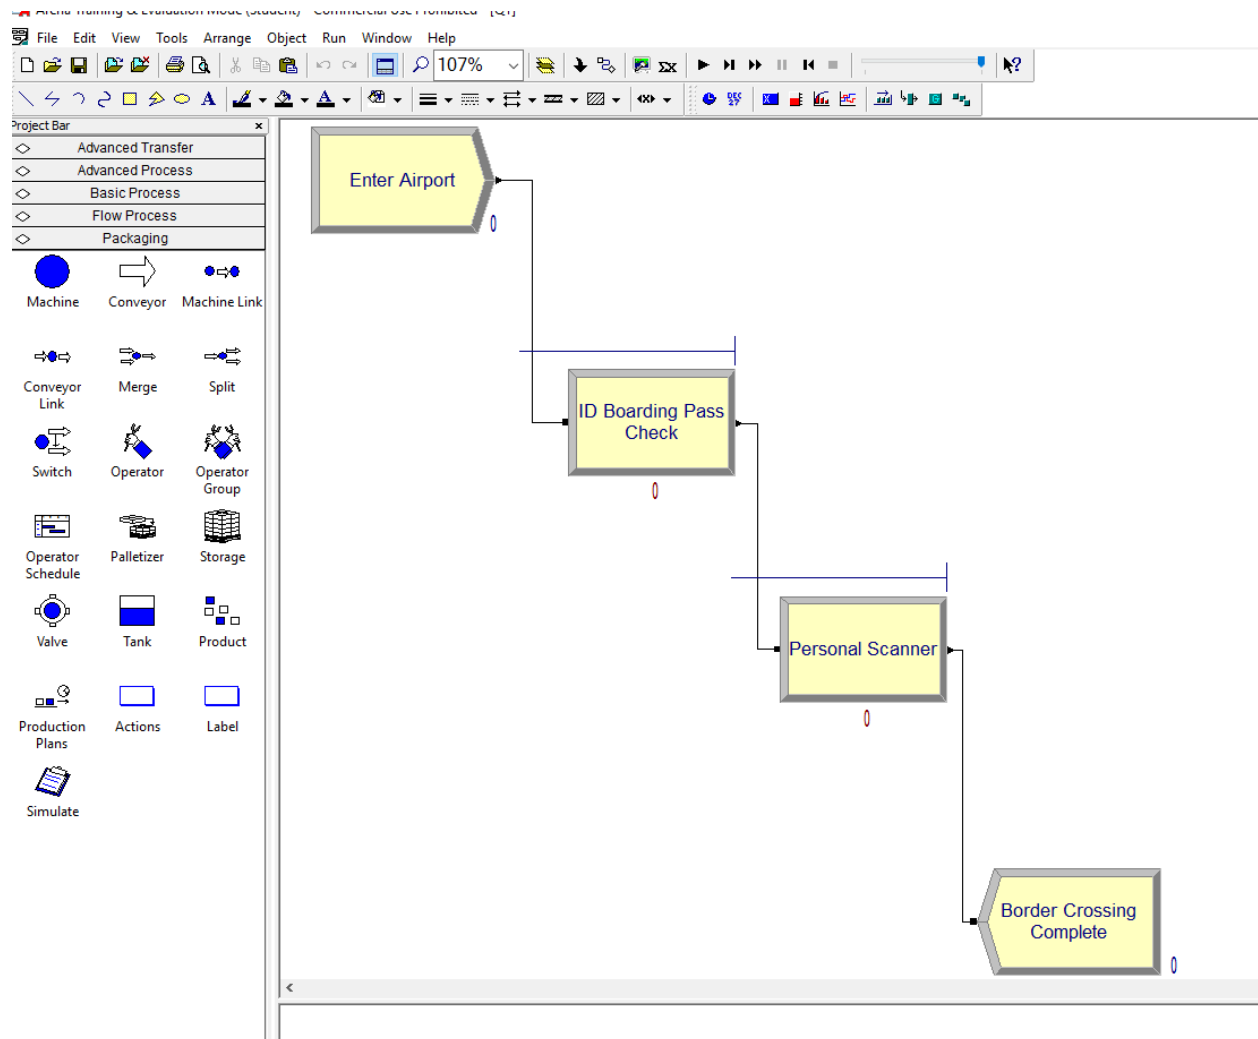


Crab

Question 1

I excuted this simulation in Arena over 10 repeats of 5 hours. It was frustrating because I hit a whole lot of errors with the simulation size. With my time over I would complete in SimPy.

Here is my basic setup:



I operated with 4 guards at the Id Check and 4 Scanners:

Set - Basic Process			
	Name	Type	Members
1 ▶	Guard Group	Resource	4 rows
2	Scanner Group	Resource	4 rows

Double-click here to add a new row.

I pretty much couldn't reduce resources or the simulation would hit the 150 error.

The results are included in the attached pdf. The end to end average processing time was 0.1184 hours or a little over 6 minutes.

Question 2

There is a reasonable amount of scripting in this answer, I feel like it could be done much more efficiently, in particular where I manual script out the models.

```
# initialize the data
cancer <- read.table('breast-cancer-wisconsin.data.txt', sep=',', header = FALSE)
# head(cancer, 10)
# apply(cancer, class)
# summary(cancer)

cancer.copy <- cancer
cancer.copy$V7[cancer.copy$V7=='?'] <- NA
cancer.copy$V7 <- as.numeric(cancer.copy$V7)
# summary(cancer.copy)

# Pattern of missing values
library('mice')
library(Hmisc)
```

```
## Loading required package: lattice
## Loading required package: survival
##
## Attaching package: 'survival'
## The following object is masked _by_ '.GlobalEnv':
##
## cancer
## Loading required package: Formula
## Loading required package: ggplot2
##
## Attaching package: 'Hmisc'
## The following objects are masked from 'package:base':
##
## format.pval, round.POSIXt, trunc.POSIXt, units
```

```
# how much data is missing
pMiss <- function(x){sum(is.na(x))/length(x)*100}
apply(cancer.copy,2,pMiss) # 2.3% of V7 missing, less than 5%
```

```
##      V1      V2      V3      V4      V5      V6      V7      V8
## 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 2.288984 0.000000
##      V9      V10     V11
## 0.000000 0.000000 0.000000
```

There is 2.3% of the 7th variable missing. We will first sub with the mean and built a model:

```
# impute mean
cancer.mean <- cancer.copy
cancer.mean$V7 <- with(cancer.mean, impute(cancer.mean$V7, mean))
# summary(cancer.mean)

# SVM Classifier
library("kernlab")
```

```
##
## Attaching package: 'kernlab'

## The following object is masked from 'package:ggplot2':
##
##      alpha
```

```
results = {}
#mean dataset
cancer.mean <- as.matrix(cancer.mean)
model.mean <- ksvm(cancer.mean[,1:10],cancer.mean[,11],type="C-svc",kernel="vanilladot",C=100,scaled=TRUE)

## Setting default kernel parameters
pred <- predict(model.mean,cancer.mean[,1:10])
acc.mean <- sum(pred == cancer.mean[,11]) / nrow(cancer.mean)
results['MEAN'] <- acc.mean
```

Next we built the data set with mode. For this one a defined the function for identifying the mode:

```
# define mode function
getmode <- function(v) {
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}

# impute mode
cancer.mode <- cancer.copy
cancer.mode$V7 <- with(cancer.mode, impute(cancer.mode$V7, fun=getmode(cancer.mode$V7)))
# summary(cancer.mode)

# mode dataset
cancer.mode <- as.matrix(cancer.mode)
model.mode <- ksvm(cancer.mode[,1:10],cancer.mode[,11],type="C-svc",kernel="vanilladot",C=100,scaled=TRUE)

## Setting default kernel parameters
pred <- predict(model.mode,cancer.mode[,1:10])
acc.mode <- sum(pred == cancer.mode[,11]) / nrow(cancer.mode)
```

```
results['MODE'] <- acc.mode
```

Here we impute values via regression and then build a model:

```
# imputation via regression
```

```
cancer.regr <- cancer.copy
```

```
imp <- mice(cancer.regr, method = "norm.predict", m = 1)
```

```
##
```

```
## iter imp variable
```

```
## 1 1 V7
```

```
## 2 1 V7
```

```
## 3 1 V7
```

```
## 4 1 V7
```

```
## 5 1 V7
```

```
cancer.regr <- complete(imp)
```

```
# regression dataset
```

```
cancer.regr <- as.matrix(cancer.regr)
```

```
model.regr <- ksvm(cancer.regr[,1:10], cancer.regr[,11], type="C-svc", kernel="vanilladot", C=100, scaled=TRUE)
```

```
## Setting default kernel parameters
```

```
pred <- predict(model.regr, cancer.regr[,1:10])
```

```
acc.regr <- sum(pred == cancer.regr[,11]) / nrow(cancer.regr)
```

```
results['REGRESSION'] <- acc.regr
```

I added some random number between -1 and 1 from uniform distribution. There are probably some arguments about the error distribution and scale but I just did this quick and dirty:

```
# add perturbation to regr
```

```
cancer.per <- cancer.regr
```

```
idx <- which(cancer.copy$V7 %in% NA)
```

```
rando <- runif(length(idx), -1, 1)
```

```
for (i in 1:length(idx)){
```

```
  cancer.per[idx[i], 7] <- cancer.per[idx[i], 7] + rando[i]
```

```
}
```

```
# use perturbation dataset
```

```
cancer.per <- as.matrix(cancer.per)
```

```
model.per <- ksvm(cancer.per[,1:10], cancer.per[,11], type="C-svc", kernel="vanilladot", C=100, scaled=TRUE)
```

```
## Setting default kernel parameters
```

```
pred <- predict(model.per, cancer.per[,1:10])
```

```
acc.per <- sum(pred == cancer.per[,11]) / nrow(cancer.per)
```

```
results['PERTURBATION'] <- acc.per
```

Here we omit all NA values and just build with the original data:

```
# omit NAs
```

```
cancer.orig <- as.matrix(cancer.copy)
```

```
model.orig <- ksvm(cancer.orig[,1:10], cancer.orig[,11], na.action = na.omit, type="C-svc", kernel="vanilladot", C=100, scaled=TRUE)
```

```
## Setting default kernel parameters
```

```
pred.orig <- predict(model.orig, cancer.orig[-idx,1:10])
```

```
acc.orig <- sum(pred.orig == cancer.orig[-idx,11]) / nrow(cancer.orig[-length(idx)])
```

```
results['REMOVED'] <- acc.orig
```

Here we add a binary value for if there is data missing:

```
# introduce binary value
cancer.bin <- cancer.copy
cancer.bin$m <- 0
for (i in 1:length(idx)){
  cancer.bin[idx[i],12]<-1
  cancer.bin[idx[i],7]<-0
}
cancer.bin <- as.matrix(cancer.bin)
model.bin <- ksvm(cancer.bin[, -11], cancer.bin[, 11], na.action=na.fail, type="C-svc", kernel="vanilladot")

## Setting default kernel parameters
pred <- predict(model.bin, cancer.bin[, -11])
acc.bin <- sum(pred == cancer.bin[, 11]) / nrow(cancer.bin)
results['BINARY'] <- acc.bin
```

And the results are:

```
results
```

```
##          MEAN          MODE  REGRESSION PERTURBATION      REMOVED
##    0.9699571    0.9685265    0.9685265    0.9685265    0.9470672
##          BINARY
##    0.9699571
```

We see because we fit straight to the data the results are similar. When we removed the data it becomes marginally less accurate. We could do smarter things with cross-val and test sets to determine how these imputation methods actually perform but this is ok for now.

Question 3

Optimisation may be used to maximise profits for a car company. They would control the make types according to competitive actions, and vehicle attributes.

We did a piece of work looking to optimise workflows through a major biometric system where the objective function was determined by operator hours required and the number of false matches and false non-matches that proceed through the system.